



# Shortest Paths

- The BFS search strategy can be used to find a shortest path from some starting vertex  $v$  to every other vertex in a connected graph
- This approach is useful when each edge is as good as any other, in particular it fails when a **weighted graph** has to be taken into account
- A **weighted graph** is a graph that has a numeric label  $w(e)$  associated with each edge  $e$  (the weight of that edge)
- For  $e = (u, v)$  we have  $w(u, v) = w(e)$

# Shortest Paths



# Shortest Paths



Possible paths from JFK to LAX:

Path	Weights	Total
JFK - ORD - SFO - LAX	$740 + 1846 + 337$	2923
JFK - ORD - DFW - LAX	$740 + 802 + 1235$	2777
JFK - MIA - LAX	$1090 + 2342$	3432
JFK - MIA - DFW - LAX	$1090 + 1121 + 1235$	3446
JFK - BOS - ORD - SFO - LAX	$187 + 867 + 1846 + 337$	3237
JFK - BOS - ORD - DFW - LAX	$187 + 867 + 802 + 1235$	3091

# Shortest Paths

The weight (or length) of a path  $P$  in a weighted graph  $G$  is the sum of the weights of the edges of  $P$ .

If  $P = ((v_0, v_1), (v_1, v_2), \dots, (v_{k-1}, v_k))$  then

$$w(P) = \sum_{i=0}^{k-1} w(v_i, v_{i+1})$$

The **distance** from a vertex  $u$  to a vertex  $v$  in  $G$ , denoted  $d(u, v)$  is the length of a minimum-length path (**shortest path**) from  $u$  to  $v$  (if such a path exists)

# Shortest Paths

## Dijkstra's Algorithm

- A class of algorithms solve the problem of finding a shortest path from some vertex  $s$  to each other vertex in a weighted graph  $G$
- One of these, **Dijkstra's algorithm**, applies the **greedy method**: a given problem is solved by repeatedly selecting the best choice from among those available at each iteration
- The main idea is to perform a weighted BFS search starting from the source vertex  $s$  and creating a “cloud” of vertices, each entering the cloud in order of their distance from  $s$
- Then, in each iteration, the next vertex chosen is the vertex outside the cloud closest to  $s$
- The algorithm terminates when no more vertices are outside the cloud (or when those outside the cloud are not connected to those within the cloud)

# Shortest Paths

## Edge Relaxation

- Let us define a label  $D[v]$  for each vertex  $v$  in  $V$
- These labels will always store the length of the best path obtained **so far** from  $s$  to  $v$
- Initially,  $D[s] = 0$  and  $D[v] = \infty$  for each  $v \neq s$
- We define the set  $C$  (the “cloud” of vertices) as the empty set
- At each iteration, a vertex  $u$  is selected not in  $C$  with smallest  $D[u]$  label, and  $u$  is put in  $C$

# Shortest Paths

## Edge Relaxation

- Once a new vertex  $u$  is pulled into  $C$ , the label  $D[v]$  of each vertex adjacent to  $u$  and outside of  $C$  is updated, to reflect the fact that may be a new and better way to reach  $v$  via  $u$
- This update is known as **relaxation**, because it takes an old estimate and checks if it can be improved to get closer to its value

More precisely:

---

```
if  $D[u] + w(u, v) < D[v]$  then  
     $D[v] = D[u] + w(u, v)$ 
```

---

# Shortest Paths

## Edge Relaxation

---

Algorithm ShortestPath( $G, s$ ):

$D[s] = 0$

$D[v] = \text{infinite}$  for each vertex  $v \neq s$

Let a priority queue  $Q$  contain all the vertices of  $G$  using the  $D$  labels as keys

while  $Q$  is not empty do

$u =$  value returned by  $Q.\text{remove\_min}()$

    for each vertex  $v$  adjacent to  $u$  such that  $v$  is in  $Q$  do

        if  $D[u] + w(u, v) < D[v]$  then

$D[v] = D[u] + w(u, v)$

            Change to  $D[v]$  the key of vertex  $v$  in  $Q$

return the label  $D[v]$  of each vertex  $v$

---



# Shortest Paths

## Computational Cost

---

. . .

Let a priority queue  $Q$  contain **all** the vertices of  $G$  using the  $D$  labels as keys

```
while  $Q$  is not empty do
     $u$  = value returned by  $Q.remove\_min()$ 
    for each vertex  $v$  adjacent to  $u$  such that  $v$  is in  $Q$ 
        do
            if  $D[u] + w(u,v) < D[v]$  then
                 $D[v] = D[u] + w(u,v)$ 
                Change to  $D[v]$  the key of vertex  $v$  in  $Q$ 
return the label  $D[v]$  of each vertex  $v$ 
```

---



# Shortest Paths

## Computational Cost


- The nested **for** loop runs in  $O(m)$  time
- The outer **while** loop executes in  $O(n)$  time
- The update of  $Q$  requires  $O(m)$  time
- If  $Q$  is a priority queue, each of the above operations run in  $O(\log n)$ , therefore the overall running time is  $O((n + m)\log n)$ , that is (approx)  $O(n^2 \log n)$

# Shortest Paths

## Dijkstra's Algorithm





## Dijkstra's Algorithm




# Shortest Paths

## Dijkstra's Algorithm





## Dijkstra's Algorithm



# Shortest Paths

## Dijkstra's Algorithm



# Dijkstra's Algorithm

## Shortest Path Tree

---


```
def shortest_path_tree(g, s, d):  
  
    tree = {}  
    for v in d:  
        if v is not s:  
            for e in g.incident_edges(v, False):  
                u = e.opposite(v)  
                wgt = e.element()  
                if d[v] == d[u] + wgt:  
                    tree[v] = e  
    return tree
```

---




# Shortest Paths

## Dijkstra's Algorithm



# Shortest Paths

## Dijkstra's Algorithm



---

```
graph = {'s': {'a': 2, 'b': 1},  
        'a': {'s': 3, 'b': 4, 'c': 8},  
        'b': {'s': 4, 'a': 2, 'd': 2},  
        'c': {'a': 2, 'd': 7, 't': 4},  
        'd': {'b': 1, 'c': 11, 't': 5},  
        't': {'c': 3, 'd': 5}}
```

---

# Shortest Paths

## Dijkstra's Algorithm

---

```
def dijkstra(graph, src, dest, visited=[], distances={},  
             predecessors={}):  
    if src not in graph:  
        raise TypeError('the root of the shortest path  
                        tree cannot be found in the graph')  
    if dest not in graph:  
        raise TypeError('the target of the shortest path  
                        cannot be found in the graph')  
    if src == dest:  
        path=[]  
        pred=dest  
        while pred != None:  
            path.append(pred)  
            pred=predecessors.get(pred, None)  
        print('shortest path: ' +str(path)+" cost="+str(  
            distances[dest]))  
    else :  
        if not visited:  
            distances[src]=0
```

# Shortest Paths

## Dijkstra's Algorithm

---

```
. . .

for neighbor in graph[src] :
    if neighbor not in visited:
        new_distance = distances[src] + graph[src][neighbor]
        if new_distance < distances.get(neighbor, float('inf'))
            :
            distances[neighbor] = new_distance
            predecessors[neighbor] = src

visited.append(src)
unvisited={}

for k in graph:
    if k not in visited:
        unvisited[k] = distances.get(k, float('inf'))
x=min(unvisited, key=unvisited.get)
dijkstra(graph,x,dest,visited,distances,predecessors)
```

---

# Minimum Spanning Trees

- All the computers in an office must be connected using the least amount of cable (wireless is not admitted)
- This problem can be modeled using an undirected weighted graph  $G$  whose vertices represent the computers, and whose edges represent all the possible pairs  $(u, v)$  of computers
- The weight  $w(u, v)$  of the edge  $(u, v)$  is equal to the amount of cable needed to connect computer  $u$  to computer  $v$
- The problem is finding a tree that contains all the vertices of  $G$  and has the minimum total weight over all such trees


# Minimum Spanning Trees

Given an undirected weighted graph  $G$ , we are interested in finding a tree  $T$  that contains all the vertices in  $G$  and minimizes the sum

$$w(T) = \sum_{(u,v) \in T} w(u, v)$$

Such a tree that contains every vertex of a connected graph  $G$  is said to be a **spanning tree**, and the problem of computing a spanning tree  $T$  with smallest total weight is known as the **minimum spanning tree (MST)** problem

# Minimum Spanning Trees



## Proposition

Let  $G$  be a weighted connected graph, and let  $V_1$  and  $V_2$  be a partition of the vertices of  $G$  into two disjoint nonempty sets. Let  $e$  be an edge in  $G$  with minimum weight from among those with one endpoint in  $V_1$  and the other  $V_2$ . There is a minimum spanning tree  $T$  that has  $e$  as one of its edges.

# Minimum Spanning Trees

## Prim-Jarnik Algorithm

- In this algorithm, a minimum spanning tree is grown starting from a single cluster starting from some “root” vertex  $s$
- The main idea is similar to that of Dijkstra's algorithm: a cloud of vertices is defined that grows at each iteration
- At each iteration, a minimum-weight edge is chosen  $e = (u, v)$  which connects a vertex  $u$  in the cloud  $C$  to a vertex  $v$  outside the cloud  $C$
- The vertex  $v$  is brought into the cloud  $C$  and the process is repeated until a spanning tree is formed
- As in Dijkstra's algorithm, a label  $D[v]$  is maintained for each vertex outside the cloud  $C$ , so that  $D[v]$  stores the weight of the minimum observed edge for joining  $v$  to the cloud  $C$



# Minimum Spanning Trees

## Prim-Jarnik Algorithm

---

Algorithm Prim\_Jarnik( $G$ ):

Input: An undirected, weighted, connected graph  $G$   
with  $n$  vertices and  $m$  edges

Output: A minimum spanning tree  $T$

Pick any vertex  $s$  of  $G$

$D[s] = 0$

for each vertex  $v \neq s$  do

$D[v] = \text{infinite}$

Initialize  $T = \emptyset$

Initialize a priority queue  $Q$  with an entry  $(D[v], (v, \text{None}))$  for each vertex  $v$ , where  $D[v]$  is the key in the priority queue, and  $(v, \text{None})$  is the associated value

while  $Q$  is not empty do

$(u, e) = \text{value returned by } Q.\text{remove\_min}()$

    Connect vertex  $u$  to  $T$  using edge  $e$

    for each edge  $e' = (u, v)$  such that  $v$  is in  $Q$  do  
        if  $w(u, v) < D[v]$  do



# Minimum Spanning Trees

## Analysis of the Prim-Jarnik Algorithm

- $n$  insertions are performed into  $Q$
- Later,  $n$  extract-min are performed
- A total of  $m$  priorities are updates
- With a priority queue, each operation runs in  $O(\log n)$ , and the overall time for the algorithms is  $O((n + m)\log n)$ , which is  $O(m\log n)$  for a connected graph
- By using an unsorted list, the running time will be  $O(n^2)$



# Minimum Spanning Trees

## Prim-Jarnik Algorithm





# Minimum Spanning Trees

## Prim-Jarnik Algorithm





# Minimum Spanning Trees

## Prim-Jarnik Algorithm





# Minimum Spanning Trees

## Prim-Jarnik Algorithm



# Minimum Spanning Trees

## Prim-Jarnik Algorithm



# Minimum Spanning Trees

## Prim-Jarnik Algorithm

---

```
def MST_PrimJarnik(g):  
    d = {}  
    tree = []  
    pq = AdaptableHeapPriorityQueue()  
    pqlocator = {}  
  
    for v in g.vertices():  
        if len(d) == 0:  
            d[v] = 0  
        else:  
            d[v] = float('inf')  
        pqlocator[v] = pq.add(d[v], (v, None))  
  
    . . .
```

(to be continued)

---



# Minimum Spanning Trees

## Prim-Jarnik Algorithm

---

. . .

```
while not pq.is_empty():
    key,value = pq.remove_min()
    u,edge = value
    del pqlocator[u]
    if edge is not None:
        tree.append(edge)
    for link in g.incident_edges(u):
        v = link.opposite(u)
        if v in pqlocator:
            wgt = link.element()
            if wgt < d[v]:
                d[v] = wgt
                pq.update(pqlocator[v], d[v], (v, link))

return tree
```

---

# Minimum Spanning Trees

## Kruskal's Algorithm

- Kruskal's algorithm maintains a **forest** of clusters, repeatedly merging pairs of clusters until a single cluster spans the graph
- Initially, each vertex is by itself a singleton cluster
- The algorithm considers each edge in turn, ordered by increasing weight
- If an edge  $e$  connects two different clusters, then  $e$  is added to the set of edges of the MST, and the clusters connected by  $e$  are merged into a single cluster
- On the contrary, if  $e$  connects two vertices that are already in the same cluster, then  $e$  is discarded
- The algorithm ends when enough edges have been added to form a spanning tree

# Minimum Spanning Trees

## Kruskal's Algorithm

---

Algorithm Kruskal( $G$ ):

Input: A simple connected weighted graph  $G$   
with  $n$  vertices and  $m$  edges

Output: A minimum spanning tree  $T$

for each vertex  $v$  in  $G$  do

    Define an elementary cluster  $C(v) = \{v\}$

    Initialize a priority queue  $Q$  to contain all edges  
    in  $G$ , using the weights as keys

$T = \emptyset$

    while  $T$  has fewer than  $n-1$  edges do

$(u,v) = \text{value returned by } Q.\text{remove\_min}()$

        Let  $C(u)$  be the cluster containing  $u$ , and let  $C(v)$   
        be the cluster containing  $v$

        if  $C(u) \neq C(v)$  then

            Add edge  $(u,v)$  to  $T$

            Merge  $C(u)$  and  $C(v)$  into one cluster

    return tree  $T$



# Minimum Spanning Trees

## Kruskal's Algorithm

- The correctness of Kruskal's algorithm is based upon the crucial fact about minimum spanning trees from Proposition
- Each time that Kruskal's algorithm adds an edge  $(u, v)$  to the MST  $T$ , we can define a partitioning of the set of vertices  $V$  by letting  $V_1$  be the cluster containing  $v$  and letting  $V_2$  contain the rest of the vertices in  $V$
- This defines a disjoint partitioning of the vertices of  $V$
- Moreover, since we are extracting edges from  $Q$  in order by their weights,  $e$  must be a minimum-weight edge with one vertex in  $V_1$  and the other in  $V_2$ .
- Therefore, Kruskal's algorithm always adds a valid MST edge



# Minimum Spanning Trees

## Kruskal's Algorithm





# Minimum Spanning Trees

## Kruskal's Algorithm





# Minimum Spanning Trees

## Kruskal's Algorithm



# Minimum Spanning Trees



## Kruskal's Algorithm







# Minimum Spanning Trees

## Kruskal's Algorithm





# Minimum Spanning Trees

## Kruskal's Algorithm



# Minimum Spanning Trees

## Kruskal's Algorithm



# Minimum Spanning Trees

## Analysis of Kruskal's Algorithm

- The ordering of edges can be implemented in  $O(m \log n)$  (by using a priority queue)
- To implement Kruskal's algorithm, we must be able to find the cluster for vertices  $u$  and  $v$  that are endpoints of an edge  $e$ , to test whether those two clusters are distinct, and if so, to merge those two cluster into one
- The management of **disjoint partitions** can be performed in  $O(m + n \log n)$
- For a connected graph we have  $m \geq n - 1$  and therefore the dominant term is  $O(m \log n)$ , which is the running time of Kruskal's algorithm

# Minimum Spanning Trees

## Kruskal's Algorithm

---

```
def MST_Kruskal(g):  
    tree = []  
    pq = HeapPriorityQueue()  
    forest = Partition()  
    position = {}  
  
    for v in g.vertices():  
        position[v] = forest.make_group(v)  
  
    for e in g.edges():  
        pq.add(e.element(), e)  
  
    size = g.vertex_count()  
    while len(tree) != size - 1 and not pq.is_empty():  
        weight, edge = pq.remove_min()  
        u, v = edge.endpoints()  
        a = forest.find(position[u])  
        b = forest.find(position[v])  
        if a != b:
```

# Minimum Spanning Trees

## Disjoint Partitions

- A partition data structure manages an universe of elements that are organized into disjoint sets (an element belongs to one and only one of these sets)
- For clarity reason, the clusters of a partition are referred to as **groups**
- To differentiate between one group and another, we assume that at any point in time, each group has a designated entry called **leader** of the group

# Minimum Spanning Trees

## Disjoint Partitions

Method	Functionality
<code>make_group(x)</code>	Create a singleton group containing new element $x$ and return the position storing $x$
<code>union(p, q)</code>	Merge the groups containing positions $p$ and $q$
<code>find(p)</code>	Return the position of the leader of the group containing position $p$